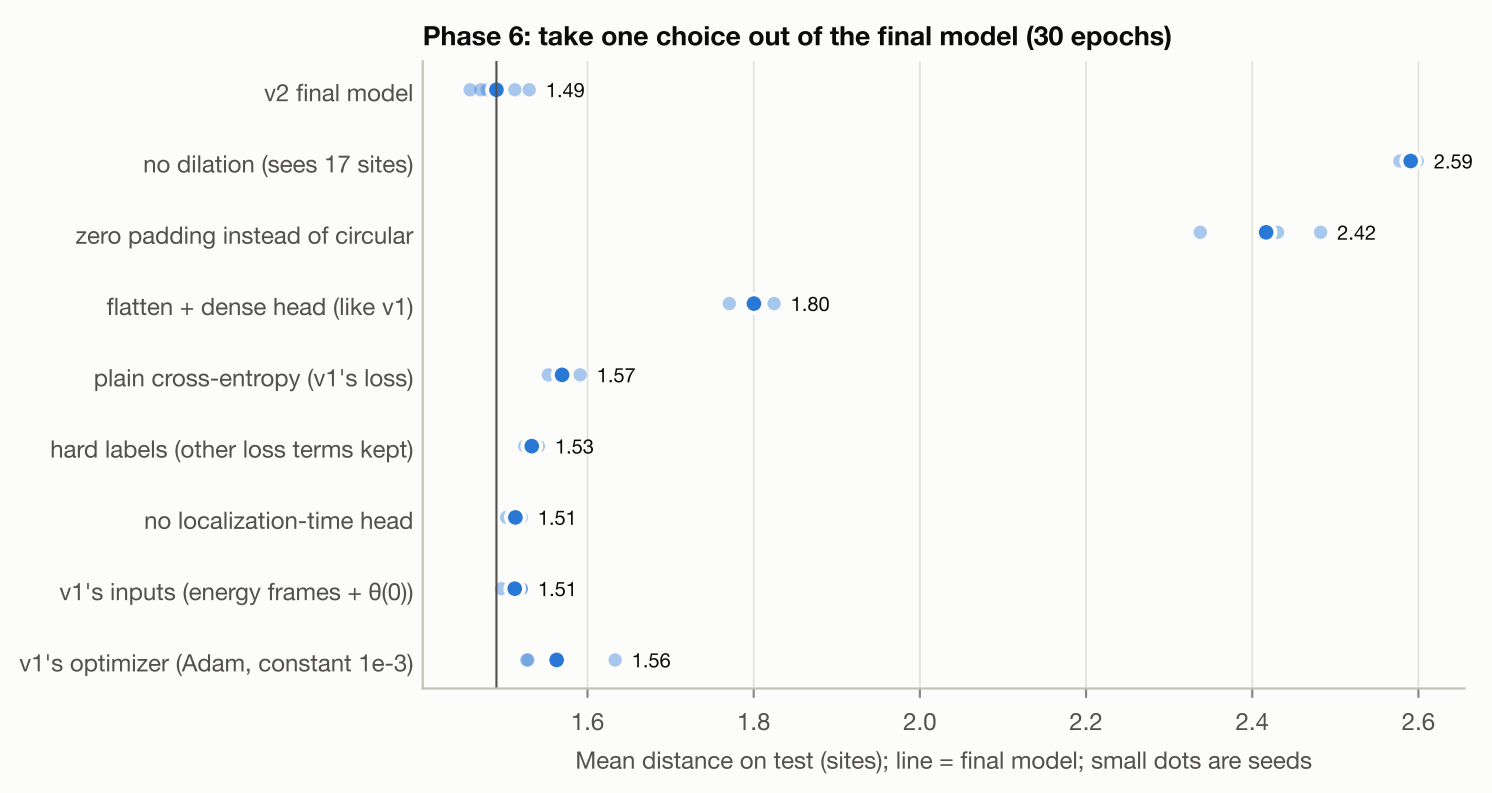
Phase 6: take one choice out of the final model (30 epochs)
v2 final model
no dilation (sees 17 sites)
zero padding instead of circular
flatten + dense head (like v1)
plain cross-entropy (v1's loss)
hard labels (other loss terms kept)
no localization-time head
v1's inputs (energy frames + θ(0))
v1's optimizer (Adam, constant 1e-3)
1.49
2.59
2.42
1.80
1.57
1.53
1.51
1.51
1.56
1.6
1.8
2.0
2.2
2.4
2.6
Mean distance on test (sites); line = final model; small dots are seeds